



BITS Pilani
Pilani Campus

Computer Networks (CS F303)

Virendra Singh Shekhawat
Department of Computer Science and Information Systems



BITS Pilani
Pilani Campus

Second Semester 2020-2021

Module-2 Application Layer

Today's Agenda



- P2P Applications
- Distributed Hash Tables
- Circular DHT Protocol
 - Chord Protocol

Peer to Peer (P2P) Architecture

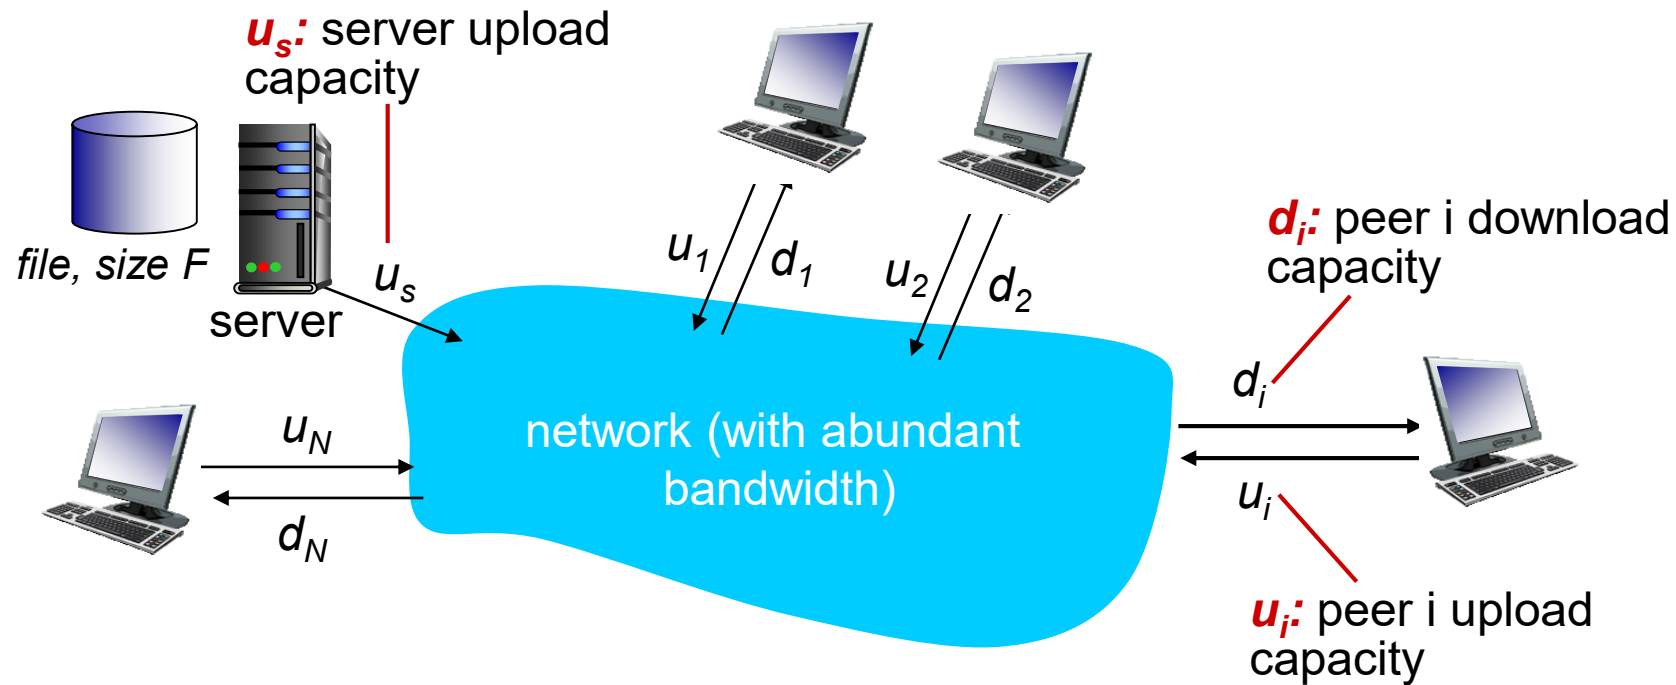
- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected
- Examples
 - File distribution (BitTorrent)
 - Streaming (KanKan)
 - VoIP (Skype)

File Distribution: P2P vs CS



Q: How much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File Distribution Time




- *Server transmission*: must sequentially send (upload) N file copies:

- Time to send N copies: NF/u_s

- ❖ *Client*: each client must download file copy

- d_{\min} = min client download rate
- Slowest client download time: F/d_{\min}



time to distribute F
to N clients using
client-server approach $D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$

File Distribution Time: P2P



- **Server transmission:** must upload at least one copy


- time to send one copy: F/u_s

- ❖ **Client:** each client must download file copy

- Slowest client download time: F/d_{\min}

- ❖ **Clients:** as aggregate must download NF bits

- max upload rate (limiting max download rate) is $u_s + \sum u_i$



Time to distribute F
to N clients using
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

Exercise




- Distributing a File $F = 15$ Gbits to 10 peers
- Server upload rate is $u_s = 30$ Mbps
- Each peer download rate is $d_i = 2$ Mbps
- Each peer upload rate is $u = 300$ Kbps
- Question
 - Calculate minimum distribution time for both CS and P2P

CS vs P2P: Example



client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$




P2P File Distribution: BitTorrent





- File divided into 64 KB to 1 MB size (typically 256 KB) chunks
- Trackers
 - Tracks peers participating in torrent
 - New peer joins torrent and registers with tracker to get list of peers, connects to subset of peers
- Torrent
 - Group of peers exchange chunks of a particular file
- Peers in torrent send/receive file chunks
 - At any given time, each peer will have a subset of chunks from the file
 - A peer asks its neighbors for the list of chunks they have and gets list from each
 - A peer needs to take a call on-
 - Which chunks should it request first from its neighbor?
 - To which of its neighbors it should send requested chunks?

The lookup problem




Decentralized network with several peers (servers/clients)
How to find specific peer that hosts desired data within this network?

P2P Protocols



Kazaa (Skype is based on Kazaa)



Distributed Database



- Each Peer hold a small subset of the total (key, value) pairs
- Any Peer can query the distributed database with a particular key
 - Distributed DB locate the Peers that have the corresponding (key, value) pairs and return to the querying Peer
 - Any Peer can insert new (key, value) pairs into the DB

Distributed Hash Table Implementation [.1]



- Randomly scatter the (key, value) pairs across all the peers
- Each peer maintain a list of the IP addresses of all peers
- The querying peer sends its query to all other peers
- The peers containing the (key, value) pairs that match the key can respond with matching pairs
- This approach is not scalable. Why?

Database Implementation [..2] Circular DHT




- Hash function assigns each “**node**” and “**key**” an m -bit *identifier* using a base hash function such as SHA-1
 - Node_ID = hash(IP, Port)
 - Key_ID = hash(original key)

ID Space: 0 to $2^m - 1$


Here: $m = 6$

Range = 64

Assign (key-value) pair to the peer that has the *closest* ID.



Chord Protocol: Lookup Operation Example



Predecessor: pointer to the previous node on the id circle

Successor: pointer to the succeeding node on the id circle


- ask node *n* to find the successor of *id*
- If *id* between *n* and its **successor** return **successor**
- else forward query to *n*'s **successor** and so on

=> #messages linear in #nodes

Scalable node localization




- Each node n contains a routing table with up-to m entries (m : number of bits of the identifier) => **finger table**
- i^{th} entry in the table at node n contains the first node s that succeeds n by at least 2^{i-1}
 - $s = \text{successor}(n + 2^{i-1})$
 - s is called the i^{th} finger of node n



The Chord algorithm – Scalable node localization



- Search in finger table for the node which is **most immediatly precedes key**
- Invoke **find_successor** from that node



Number of messages $O(\log N)$!

Failure Recovery (Peer Churn)

- Key step in failure recovery is maintaining correct successor pointers
- To achieve this, *each node maintains a successor-list* of its r nearest successors on the ring
- *If node n notices that its successor has failed*, it replaces it with the first live entry in the list
- The ***stabilize*** will correct finger table entries and successor-list entries pointing to failed node
- Stabilization protocol should be invoked based on the frequency of nodes leaving and joining

Thank You!